

## ARTIFICIAL INTELLIGENCE PyQ ANSWERS

### Q. Define the relationship between agent and environment

- **Agent:**  
An **agent** is any entity that perceives its environment through sensors and acts upon that environment through actuators. It could be a robot, a software bot, a self-driving car, or even an AI model in a game.
- **Environment:**  
The **environment** is the external world in which the agent operates. It provides the context and conditions for the agent's actions, and it responds to those actions, sometimes changing as a result.

### Agent-Environment Relationship

#### 1. Perception-Action Cycle

The interaction between the agent and environment occurs through a continuous cycle:

1. **Perceive:**  
The agent uses its sensors to **observe the environment** (e.g., a camera, temperature sensor, or data stream).
2. **Decide:**  
Based on its observations and internal logic or model, the agent **makes decisions**.
3. **Act:**  
The agent then uses its actuators to **take actions** that affect the environment.
4. **Feedback:**  
The environment changes based on the agent's actions, and the agent perceives the updated environment in the next cycle.

**This loop is often called the Sense-Think-Act cycle.**

### Types of Environments

The nature of the environment affects how the agent behaves:

Environment Type	Description
<b>Fully Observable</b>	Agent has complete information (e.g., chess).
<b>Partially Observable</b>	Agent sees only part of the environment (e.g., driving in fog).
<b>Deterministic</b>	Outcomes are predictable (e.g., mathematical calculations).
<b>Stochastic</b>	Outcomes are probabilistic (e.g., poker).
<b>Static</b>	Environment doesn't change while the agent is thinking (e.g., crossword puzzle).
<b>Dynamic</b>	Environment changes over time (e.g., real-time video game).
<b>Discrete</b>	Finite number of states/actions (e.g., board games).
<b>Continuous</b>	Infinite possibilities (e.g., driving).

### Goal of the Agent

The goal of an agent is to:

- Maximize **performance** based on a performance measure.
- Make decisions that lead to the **best possible outcome** in its environment over time.

**Q. Name the different types of environments and briefly explain effects of each environment on agent.**

**1. Fully Observable Environment**

- **Definition:** The agent has access to the complete state of the environment at every point in time.
- **Example:** Chess, where all pieces and their positions are visible.
- **Effect on Agent:**
  - Easier to design optimal agents.
  - The agent can make decisions based on full knowledge.
  - Reduces uncertainty in decision-making.

**2. Partially Observable Environment**

- **Definition:** The agent has limited or incomplete access to the environment's state.
- **Example:** Poker, where opponents' cards are hidden.
- **Effect on Agent:**
  - The agent must maintain internal states or beliefs (memory/history).
  - Decision-making is probabilistic and based on assumptions or estimations.
  - Increases complexity and uncertainty.

**3. Deterministic Environment**

- **Definition:** The next state of the environment is completely determined by the current state and the agent's action.
- **Example:** A puzzle where each move leads to a predictable result.
- **Effect on Agent:**
  - Easier to plan and predict outcomes.
  - Allows for precise strategies without randomness.
  - No need for probabilistic reasoning.

**4. Stochastic Environment**

- **Definition:** The next state of the environment is influenced by randomness or uncertainty.
- **Example:** Driving a car on a road with unpredictable pedestrian movements.
- **Effect on Agent:**
  - The agent must account for probabilities.
  - Strategies must be flexible and adaptive.
  - Requires robust handling of uncertainty.

**5. Episodic Environment**

- **Definition:** The agent's experience is divided into distinct episodes, and each decision does not depend on previous actions.
- **Example:** Image recognition tasks where each image is independent.
- **Effect on Agent:**
  - Simplifies the agent's decision-making.
  - No need for long-term planning or memory.
  - Actions are evaluated in isolation.

**6. Sequential Environment**

- **Definition:** Current decisions affect future outcomes; episodes are interconnected.
- **Example:** Chess or navigation tasks.
- **Effect on Agent:**
  - The agent must plan ahead and consider future consequences.

- Requires memory and foresight.
- More complex decision-making process.

## 7. Static Environment

- **Definition:** The environment remains unchanged while the agent is making decisions.
- **Example:** Crossword puzzles.
- **Effect on Agent:**
  - No need to worry about time constraints.
  - Easier to deliberate and optimize decisions.
  - Agent can use search and planning techniques effectively.

## 8. Dynamic Environment

- **Definition:** The environment changes over time, possibly independent of the agent's actions.
- **Example:** Autonomous driving, where other vehicles move independently.
- **Effect on Agent:**
  - The agent must act quickly and adapt continuously.
  - Time-sensitive decision-making is critical.
  - May require real-time data processing.

## 9. Discrete Environment


- **Definition:** The environment consists of a finite number of distinct states, actions, and events.
- **Example:** Board games.
- **Effect on Agent:**
  - Easier to model using rules or tables.
  - Ideal for logic-based and rule-based systems.
  - Suitable for classical AI algorithms.

## 10. Continuous Environment

- **Definition:** The environment has a range of possible values for states and actions.
- **Example:** Robot arm movements or flying a drone.
- **Effect on Agent:**
  - Requires techniques from continuous mathematics (e.g., calculus, control theory).
  - Often involves approximation and learning-based strategies.
  - Harder to model exhaustively.

Each environment type defines **constraints and capabilities** for the agent. Understanding the nature of the environment helps in designing better agents and choosing suitable AI techniques.

Q. In a grid world environment, the goal of the agent is to reach state I starting from state A without visiting the shaded states. In each of the states, the agent can perform any of the four actions: up, down, left, and right to achieve the goal.



Explain what is the outcome of stochastic policy with help of the grid world environment, assuming given a state A, and suppose the stochastic policy returns the probability distribution over the action space as [0.10, 0.70, 0.10, 0.10].

In a **stochastic policy**, instead of choosing a fixed action in a state, the agent chooses actions **based on a probability distribution** over the possible actions. Let's apply this concept to the grid world environment shown in your image.

**Given:**

- The agent starts at **state A**.
- The **goal** is to reach **state I**.
- The **shaded cells (B, C, G, H)** are **not allowed** (obstacles).
- The **action space** = [Up, Down, Left, Right]
- The stochastic policy for state A is:  
 $P(\text{actions}) = [\text{Up}=0.10, \text{Down}=0.70, \text{Left}=0.10, \text{Right}=0.10]$

**Interpreting the Policy:**

The policy is:




- 10% chance to go **Up**
- 70% chance to go **Down**
- 10% chance to go **Left**
- 10% chance to go **Right**

In **state A**, the only valid actions are:

- **Down** → moves to state **D**
- **Right** → moves to **B** (✗ invalid, shaded)
- **Up** and **Left** → go outside the grid (✗ invalid)

**Outcome of the Policy in Practice:**

Since this is a **stochastic policy**, each time the agent is in state A, it samples an action based on the given probabilities.

- **70% of the time**, the agent goes **Down** to **state D** →  valid move
- **10% of the time**, it tries to go **Right** to **state B** →  invalid (blocked)
- **10% for Up** and **10% for Left** →  invalid (outside the grid)

So, **only 70% of the time**, the agent will make a successful move from A to D.

In the other **30%** of the times, it attempts an invalid move and likely **remains in state A**.

#### What This Means for the Agent:

- **Progress is slower** than a deterministic policy that always chooses "Down."
- **Uncertainty** in action selection can cause the agent to **take longer paths** or **get stuck** if not properly handled.
- With enough time (iterations), it can still reach the goal due to the probabilistic nature allowing for correct moves.

#### Conclusion:





A **stochastic policy** introduces randomness. Even if the optimal action is known (e.g., going "Down" in A), there's a chance the agent will try less optimal or even invalid moves. In your grid world, this means slower or inefficient paths, but it's essential in many RL problems to encourage **exploration**.

#### Effect of Stochasticity:

Because of randomness:

- At **A**, there's a 30% chance the agent chooses an invalid move and stays in **A**.
- Similar risks apply in **D**, **E**, and **F** if the agent doesn't pick the optimal action.
- Thus, the agent **may take longer** to reach I, or **may oscillate** until it learns or randomly chooses the correct sequence.

#### Summary:

Step	State	Action Taken	Success?
1	A	Down (70%)	 Move to D
2	D	Right	 Move to E
3	E	Right	 Move to F
4	F	Down	 Move to I (Goal)

## Q. How the Q-function differs from the value function of Reinforcement Learning?

### Value Function (V-function)

- **Denoted as:**  $V(s)$
- **Definition:** The **value function** estimates the expected return (future cumulative reward) starting from state  $s$ , and following a certain policy  $\pi$ .
- **Formula:**

$$V^{\pi}(s) = \mathbb{E}_{\pi} \left[ \sum_{t=0}^{\infty} \gamma^t R_{t+1} \mid S_0 = s \right]$$

- **Interpretation:** How good it is to be in state  $s$ , assuming the agent follows policy  $\pi$ .

### Q-Function (Action-Value Function)

- **Denoted as:**  $Q(s, a)$
- **Definition:** The **Q-function** estimates the expected return starting from state  $s$ , taking action  $a$ , and thereafter following policy  $\pi$ .
- **Formula:**

$$Q^{\pi}(s, a) = \mathbb{E}_{\pi} \left[ \sum_{t=0}^{\infty} \gamma^t R_{t+1} \mid S_0 = s, A_0 = a \right]$$

- **Interpretation:** How good it is to take action  $a$  in state  $s$ , assuming the agent follows policy  $\pi$  afterward.

Aspect	Value Function V(s)	Q-Function Q(s, a)
<b>Definition</b>	Expected cumulative reward starting from state s, following a policy $\pi$ .	Expected cumulative reward starting from state s, taking action a, then following policy $\pi$ .
<b>Formula</b>	$V^\pi(s) = \mathbb{E}_\pi \left[ \sum_{t=0}^{\infty} \gamma^t R_{t+1} \mid S_0 = s \right]$	$Q^\pi(s, a) = \mathbb{E}_\pi \left[ \sum_{t=0}^{\infty} \gamma^t R_{t+1} \mid S_0 = s, A_0 = a \right]$
<b>Meaning</b>	Expected reward from a state	Expected reward from a state-action pair
<b>Used for</b>	Evaluating states	Evaluating actions
<b>What it Evaluates</b>	How good it is to be in a given state under policy $\pi$	How good it is to take a particular action in a given state under policy $\pi$
<b>Use Case</b>	Evaluates <b>states</b>	Evaluates <b>state-action pairs</b>
<b>Action Selection</b>	Needs policy to select action	Can directly choose best action via $\text{argmax}_a Q(s, a)$
<b>Policy Improvement</b>	Indirect: needs to combine with a policy (e.g., in Actor-Critic methods)	Direct: optimal policy can be obtained by maximizing Q-values
<b>Computation</b>	Aggregates over all actions under a policy	Evaluates specific actions
<b>Commonly Used In</b>	<ul style="list-style-type: none"> <li>- Policy Evaluation</li> <li>- Value Iteration</li> <li>- Actor-Critic Algorithms</li> </ul>	<ul style="list-style-type: none"> <li>- Q-Learning</li> <li>- Deep Q-Networks (DQN)</li> <li>- SARSA</li> </ul>
<b>Complexity (Dimensionality)</b>	Lower dimensionality (since it depends only on state)	Higher dimensionality (depends on both state and action)
<b>Storage Requirements</b>	Requires storing value for each state	Requires storing value for each state-action pair
<b>Learning Focus</b>	Learn <b>value of being</b> in a state	Learn <b>value of doing</b> an action in a state
<b>Suitability</b>	Suitable for environments where policy is fixed and known	Suitable for learning <b>optimal policies</b> , especially in model-free settings

**Q. Write the bellman equation of calculating updated Q-function considering state S and action A**

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ R_{t+1} + \gamma \cdot \max_{a'} Q(S_{t+1}, a') - Q(s, a) \right]$$

Where:

- $s$ : current state
- $a$ : current action
- $S_{t+1}$ : next state
- $A_{t+1} \sim \pi(\cdot | S_{t+1})$ : next action sampled from policy
- $R_{t+1}$ : reward received after taking action  $a$  in state  $s$
- $\gamma$ : discount factor ( $0 \leq \gamma \leq 1$ )
- $\alpha$ : learning rate ( $0 < \alpha \leq 1$ )
- $\max_{a'} Q(S_{t+1}, a')$ : maximum expected future reward from next state

## Q. Sussman Anomaly

The **Sussman Anomaly** is a classic problem in the field of **Artificial Intelligence (AI) planning**, particularly in **non-linear planning** or **partial-order planning**. It highlights a limitation of early planning systems that used simple goal decomposition.

Let me break it down in a simple and detailed way.

### What Is Planning in AI?

In AI, **planning** refers to generating a sequence of actions that an agent can perform to achieve a goal from an initial state.

A planning system tries to:

- Break the goal into subgoals.
- Solve each subgoal.
- Combine those solutions into a plan.

### What Is the Sussman Anomaly?

The **Sussman Anomaly** is a **counterexample** that shows how naive **goal decomposition** can fail.

It was introduced by **Gerald Jay Sussman** in the early 1970s to show that solving subgoals independently and then merging them can sometimes **fail to find a solution**.

### Problem Setup (Block World Example)

We are working in the **blocks world**, a standard toy domain in AI.

**Initial State:**

Three blocks: **A**, **B**, and **C**

ON(A, TABLE)



ON(B, TABLE)

ON(C, A)

This means:

- Block A and Block B are on the table.
- Block C is on top of Block A.

**Goal:**

ON(A, B)

ON(B, C)

This means:

- A should be on B.
- B should be on C.

### Where the Problem Occurs

If we try to solve each goal independently:

**Subgoal 1: Make ON(A, B)**

- We might stack A on B.  Done.

**Subgoal 2: Make ON(B, C)**

- Now we try to put B on C.
- But **A is on B**, so B can't be moved unless A is removed.
- We have to **undo** the previous subgoal to achieve the next one.

**So solving goals independently creates interference.**

### Why It's an "Anomaly"

In early planners (like STRIPS), goals were solved **one at a time**, in order, without considering interactions. They assumed:

- Once a subgoal is done, it stays done.

But in this case:

- Solving one subgoal prevents the next.
- Solving the second subgoal undoes the first.

This is the **Sussman Anomaly**: naive decomposition fails due to **interactions between subgoals**.

### Solution: Nonlinear or Partial-Order Planning

To solve the anomaly, we use more sophisticated planning methods that:

- Allow actions to be partially ordered.
- Interleave actions for different subgoals.
- Consider interactions (threats, dependencies).

One such planner is **NOAH (Nets of Action Hierarchies)** or later **TWEAK, SNLP**, etc.

A valid plan might be:

1. Move C to table.
2. Move B onto C.
3. Move A onto B.

Now the goal is reached **without interference**.

### Summary

Concept	Description
<b>Sussman Anomaly</b>	A problem where independent goal solving fails due to interference.
<b>Domain</b>	Blocks world (stacking blocks).
<b>Issue</b>	Solving subgoal A-B prevents solving B-C.

<b>Lesson</b>	Planning must consider <b>goal interactions</b> .
<b>Fix</b>	Use <b>non-linear planning</b> that can handle interleaved goals.

## Q. What is Markov decision process?

### Definition:

A **Markov Decision Process (MDP)** is a **mathematical framework** used to model decision-making in situations where outcomes are partly **random** and partly under the control of a **decision maker** (agent).

An MDP provides a formalization for **reinforcement learning problems** and defines the environment in which an agent operates.

### Components of MDP:

An MDP is defined by a **5-tuple (S, A, P, R, γ)**:

<b>Component</b>	<b>Description</b>
<b>S</b>	A finite set of <b>states</b> representing all possible situations the agent can be in.
<b>A</b>	A finite set of <b>actions</b> the agent can take. The set of actions may depend on the current state.
<b>P</b>	A <b>transition probability function</b> : $P(s' s,a)$ , which defines the probability of moving to state $s'$ when action $a$ is taken in state $s$ .
<b>R</b>	A <b>reward function</b> : $R(s,a,s')$ , which gives the expected reward received after transitioning from state $s$ to state $s'$ .
<b>γ</b>	A <b>discount factor</b> $0 \leq \gamma \leq 1$ , which determines the importance of future rewards. A value close to 0 makes the agent far-sighted.

### Markov Property:

The core assumption in an MDP is the **Markov Property**, which states that:

**The future state depends only on the current state and action, and not on the sequence of events that preceded it.**

Formally:

$$P(S_{t+1} \mid S_t, A_t, S_{t-1}, A_{t-1}, \dots, S_0, A_0) = P(S_{t+1} \mid S_t, A_t)$$

This "memoryless" property is essential for simplifying the problem and applying dynamic programming techniques.

### Objective of an Agent in an MDP:

The agent's goal is to learn a **policy**

$\pi(a|s)$  that **maximizes the expected cumulative reward** over time:

$$\text{Maximize } \mathbb{E}_{\pi} \left[ \sum_{t=0}^{\infty} \gamma^t R_{t+1} \right]$$

## Q. Write Q-learning algorithm

**Q-learning** is a value-based **off-policy** algorithm used to find the **optimal action-selection policy** for any given finite Markov Decision Process (MDP).

It learns the **optimal Q-values**: the expected future rewards for a state–action pair under the best policy.

Update the Q-value using the Bellman Optimality Equation:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ r + \gamma \cdot \max_{a'} Q(s', a') - Q(s, a) \right]$$

Where:

- $s$ : current state
- $a$ : action taken
- $r$ : reward received
- $s'$ : next state
- $a'$ : possible next actions
- $\alpha$ : learning rate ( $0 < \alpha \leq 1$ )
- $\gamma$ : discount factor ( $0 \leq \gamma \leq 1$ )

## Q-Learning Algorithm Steps

Step	Description
1	Initialize Q-values: $Q(s,a) \leftarrow 0$ for all state–action pairs $(s, a)$
2	Repeat for each episode (until convergence or max episodes):
3	→ Initialize the starting state $s$
4	→ Repeat (for each step in the episode):
5	→→ Choose an action $a$ from state $s$ using a <b>policy</b> (e.g., $\epsilon$ -greedy)
6	→→ Take action $a$ , observe reward $r$ and next state $s'$
7	→→ Update Q-value using:
	$Q(s, a) \leftarrow Q(s, a) + \alpha \cdot [r + \gamma \cdot \max_{a'} Q(s', a') - Q(s, a)]$
8	→→ Set $s \leftarrow s'$
9	→ Until state $s$ is terminal
10	End Episode
11	After all episodes, the optimal policy is: $\pi(s) = \arg \max_a Q(s, a)$

## ARTIFICIAL INTELLIGENCE PyQ ANSWERS ~ FROM MIDSEM

**Q. Explain the role of discount factor in RL, considering  $\gamma = 0, 1$  and varies b/w 0.2 to 0.8**

The **discount factor** ( $\gamma$ ) in Reinforcement Learning (RL) plays a crucial role in determining how much future rewards contribute to the agent's decision-making. It is a value between **0 and 1** that balances **immediate vs. future rewards** in the **return (cumulative reward)** calculation.

### 1. Mathematical Role of Discount Factor

The **return** at time step  $t$ :

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$$

where:

- $R_{t+1}$  is the reward at time  $t + 1$ ,
- $\gamma$  determines how much the agent values future rewards.

A **higher  $\gamma$**  makes the agent **more future-focused**, while a **lower  $\gamma$**  makes it **short-sighted**.

### 2. Effects of Different Values of $\gamma$

#### (i) When $\gamma = 0$

- The agent **only considers immediate rewards** and **completely ignores future rewards**.
- The return simplifies to  $G_t = R_{t+1}$ , meaning it behaves in a **greedy manner**, maximizing only the next reward.
- Useful in **one-step decision-making** problems where only the next action matters (e.g., reflex-based tasks).

#### (ii) When $\gamma = 1$

- The agent **considers the entire future rewards without discounting**.
- It aims for the **longest-term reward maximization**, making it **highly strategic**.
- However, in **infinite-horizon problems**, the return may **not converge**, making it computationally unstable.

#### (iii) When $\gamma$ varies (0.2 to 0.8)

- $\gamma = 0.2 \rightarrow$  The agent values **immediate rewards much more** and slightly considers future rewards.
- $\gamma = 0.5 \rightarrow$  The agent balances **short-term and long-term rewards**.
- $\gamma = 0.8 \rightarrow$  The agent **strongly considers future rewards**, optimizing for a **longer horizon** while still discounting somewhat.

### 3. Practical Considerations

- **Small  $\gamma$  (e.g., 0.2 - 0.4):** Good for **short-term tasks** like **robotic arms**, where immediate feedback is crucial.
- **Medium  $\gamma$  (e.g., 0.5 - 0.7):** Balanced strategy for **episodic tasks** like **board games**, where future outcomes matter but immediate actions are still important.

- **High  $\gamma$  (e.g., 0.8 - 0.99):** Preferred for **long-horizon tasks** like **autonomous driving** or **stock trading**, where long-term success is critical.

#### 4. Choosing $\gamma$ in Practice

- If  $\gamma$  is **too low**, the agent acts **short-sighted** and may **miss optimal strategies**.
- If  $\gamma$  is **too high**, the agent may **struggle with long-term credit assignment** and may not learn efficiently.

### Q. Write the properties of MINIMAX game search algorithm

The **Minimax algorithm** is used in **two-player, zero-sum games** like Chess, Tic-Tac-Toe, and Checkers. It systematically explores possible moves, assuming both players play optimally.

#### 1. Completeness

**Minimax is complete** if the search tree is finite.

It guarantees finding a solution if a terminal state exists.

**Example:** In **Tic-Tac-Toe**, Minimax explores all possible moves, ensuring it finds a winning or drawing strategy.

#### 2. Optimality

**Minimax is optimal** if both players play perfectly.

It assumes the opponent plays optimally and chooses the best possible move to minimize the worst-case loss.

**Example:** In **Chess**, Minimax ensures the best possible outcome based on available information.

#### 3. Time Complexity

**Minimax has exponential time complexity of  $O(b^d)$** , where:

- $b$  = branching factor (average number of moves per turn).
- $d$  = depth of the game tree.

**Example:** In Chess ( $b \approx 35$ ,  $d \approx 100$ ), Minimax becomes infeasible without optimizations like Alpha-Beta Pruning.

#### 4. Space Complexity

**Depends on the implementation:**

**DFS-based Minimax**  $\rightarrow O(d)$  (depth-first, stores only one path at a time).

**BFS-based Minimax**  $\rightarrow O(b^d)$  (breadth-first, stores the entire tree).

**Example:**

- **Tic-Tac-Toe** (small tree)  $\rightarrow$  Can store the full tree.
- **Chess** (huge tree)  $\rightarrow$  Uses depth-limited search with Alpha-Beta pruning.

#### 5. Deterministic & Zero-Sum

**Minimax works in deterministic, zero-sum games.**

- **Deterministic:** No randomness; every move leads to a known state.

- **Zero-Sum:** One player's gain is another's loss.

**Example:**

- **Applicable**  $\rightarrow$  Chess, Tic-Tac-Toe (Fixed moves, no randomness).
- **Not applicable**  $\rightarrow$  Poker (Random cards, bluffing).

#### 6. Limited by Depth and Pruning

**Minimax is inefficient for large trees** but can be improved using:

- **Depth-limited Minimax** – Stops at a fixed depth ( $d$ ).
- **Alpha-Beta Pruning** – Reduces explored nodes, improving efficiency.

**Example:** In Chess, Alpha-Beta Pruning reduces  $b^d$  complexity to  $b^{(d/2)}$ , making deeper searches feasible.

## Q. When do you apply Alpha-Beta Pruning in the Minimax Tree?

**Alpha-Beta Pruning** is applied when we can **avoid evaluating parts of the Minimax tree** that won't affect the final decision. It helps **reduce the number of nodes explored**, making Minimax faster **without changing the result**.

### 1. Pruning Condition

**Prune a branch if we find that it cannot influence the final decision.**

- $\alpha$  (alpha) → Best value for MAX (maximizing player) so far
- $\beta$  (beta) → Best value for MIN (minimizing player) so far
- If  $\alpha \geq \beta$ , further exploration is **useless**, and we prune that branch.

### 2. When to Apply Alpha-Beta Pruning?

**Pruning occurs in two cases:**

#### 1. Beta Cutoff ( $\beta \leq \alpha$ ) in the Maximizing Level

- If a MAX node finds a move with a value  $\geq \beta$ , further children **are ignored**.

#### 2. Alpha Cutoff ( $\alpha \geq \beta$ ) in the Minimizing Level

- If a MIN node finds a move with a value  $\leq \alpha$ , further children **are ignored**.

### 3. Benefits of Alpha-Beta Pruning

- Reduces nodes explored from  $O(b^d)$  to  $O(b^{(d/2)})$  → Much faster!
- Works best when the tree is sorted (Best moves first).
- No change in the final Minimax decision.

### 4. When to Avoid Alpha-Beta Pruning?

- If the tree is unstructured/random, pruning may not help much.
- Not useful in non-deterministic games (like Poker, where chance affects outcomes).
- Sorting moves before searching increases efficiency but adds extra cost.

## Q. What is the purpose of a Belief Network ?

A **Belief Network**, also known as a **Bayesian Network (BN)**, is a **probabilistic graphical model** that represents **dependencies** among random variables using **Directed Acyclic Graphs (DAGs)**. It is used for **reasoning under uncertainty** in AI.

### Purpose of a Belief Network

#### 1. Probabilistic Reasoning

- Helps **infer hidden (unknown) variables** based on known evidence.
- Computes the **probability of events** occurring.

**Example:**

- **Medical Diagnosis:** If a patient has a fever, what is the probability they have the flu?
- **Spam Detection:** Given features like sender and keywords, what is the probability an email is spam?

## 2. Handling Uncertainty in AI

- Real-world AI applications involve uncertainty (e.g., noisy data, incomplete info).
- Bayesian Networks model relationships probabilistically, unlike deterministic logic.

### Example:

- A **robotic system** must decide if an object in front is a wall or a door based on noisy sensor data.

## 3. Causal Relationship Representation

- Unlike simple probability models, BNs **represent cause-and-effect relationships**.
- Helps AI **predict outcomes** when conditions change.

### Example:

- **Traffic Prediction:** If it rains, what is the probability of a traffic jam?
  - Rain → Slippery Roads → More Accidents → Traffic Jam

## 4. Decision Making

- Used in **decision-making systems** to evaluate different actions and their probabilities.
- Supports **decision trees, reinforcement learning, and AI agents**.

### Example:

- **Self-driving cars:** Given current road conditions and pedestrian movement, what is the best driving action to take?

## 5. Learning from Data

- Bayesian Networks can be **built from data** using **Bayesian inference**.
- Allows AI to **learn probabilistic dependencies** and improve over time.

### Example:

- AI learns **which symptoms are highly correlated** with specific diseases by analyzing **medical datasets**.

## Q. Why is Probabilistic Reasoning Needed in AI?

### What is Probabilistic Reasoning?

Probabilistic reasoning in AI deals with **uncertainty** by assigning probabilities to different outcomes. Instead of making rigid, deterministic decisions, AI can **infer and predict outcomes based on likelihoods**. It is used in **Bayesian networks, Hidden Markov Models, Decision Trees, and Reinforcement Learning**.

## Why Do We Need Probabilistic Reasoning in AI?

### 1. Handling Uncertainty

- **Real-world data is incomplete, noisy, or ambiguous**—probabilistic reasoning allows AI to make the **best possible decision** even when full information isn't available.
- AI must **infer missing details** instead of assuming absolute truths.

### Example:

- A self-driving car detects a blurry object ahead. Is it a **pedestrian or just a shadow**?
- Using probabilities, the AI can determine the **most likely scenario** and react accordingly.

### 2. Making Rational Decisions

- AI applications like **medical diagnosis, stock prediction, and robotics** require **decision-making under uncertainty**.

- Probabilistic models help AI **weigh different possibilities** and choose the **most rational action**.

**Example:**

- In **medical diagnosis**, if a patient has symptoms A, B, and C, what is the probability of **Disease X vs. Disease Y**?
- AI computes probabilities and recommends **the most likely diagnosis**.

### 3. Learning from Data

- AI can learn **patterns and trends from data** using **probability distributions**.  
- Unlike rule-based systems, probabilistic models can **adapt and update** based on new information.

**Example:**

- A **spam filter** assigns a probability score based on **words, sender, and email history** to decide if a message is spam or not.
- If a user marks an email as spam, the AI **updates its probability model** to improve future predictions.

### 4. Modeling Cause-and-Effect

- Probabilistic reasoning allows AI to **understand causal relationships** rather than just correlations.  
- Helps in **predictive modeling** where past events influence future outcomes.

**Example:**

- **Traffic Prediction System:**
  - If it **rains**, the probability of **traffic congestion** increases.
  - If it **rains and there's an accident**, congestion probability is **even higher**.

### 5. Optimizing AI Performance

- AI models like **Hidden Markov Models (HMMs), Bayesian Networks, and Reinforcement Learning** use probability to **balance exploration vs. exploitation**.  
- This improves AI's ability to **adapt dynamically**.

**Example:**

- **Reinforcement Learning in Games**
  - AI **chooses moves based on the probability** of winning.
  - Over time, it learns which actions are more **rewarding** and adjusts its strategy.

### Applications of Probabilistic Reasoning in AI

- **Robotics** – Navigate uncertain environments.  
- **Natural Language Processing (NLP)** – Understand speech and text ambiguities.  
- **Medical Diagnosis** – Predict diseases based on symptoms.  
- **Fraud Detection** – Identify suspicious transactions using probability.  
- **Self-Driving Cars** – Make safe driving decisions under uncertainty.  
- **Weather Forecasting** – Predict rain, storms, or temperature changes.



## Difference Between Games and Search Problems in AI

Aspect	Games in AI	Search Problems in AI
Definition	Games involve <b>two or more agents</b> competing to achieve a goal, where each agent's actions affect the others.	Search problems involve finding a <b>sequence of actions</b> that leads to a desired goal state.
Number of Agents	<b>Multi-agent</b> environment with competing entities.	<b>Single-agent</b> environment, solving a problem independently.
Nature of Environment	<b>Adversarial</b> , as agents have <b>conflicting</b> objectives (e.g., one wins, the other loses).	<b>Non-adversarial</b> , as there is no competition, only finding an optimal solution.
Objective	To <b>maximize an agent's utility</b> while minimizing the opponent's success.	To find the <b>best or optimal path</b> from an initial state to the goal state.
Decision Process	Agents make <b>strategic decisions</b> based on the opponent's possible moves.	The search algorithm explores possible paths systematically to find a solution.
Types of Problems	Chess, Tic-Tac-Toe, Go, Poker, AlphaGo.	Route Planning, Puzzle Solving, Pathfinding, AI Planning.
Evaluation	Uses a <b>utility function</b> or evaluation function to decide the best move.	Uses <b>heuristics, cost functions, and goal tests</b> to evaluate paths.
Complexity	Often <b>more complex</b> due to the need to predict an opponent's moves (e.g., exponential growth in possibilities).	Complexity depends on the <b>state space</b> and branching factor but is usually <b>deterministic</b> .
Algorithms Used	<b>Minimax, Alpha-Beta Pruning, Monte Carlo Tree Search (MCTS).</b>	<b>A*, BFS, DFS, Dijkstra's Algorithm, Greedy Best-First Search.</b>
Example of States	In Chess, a state represents <b>board positions of all pieces</b> and the turn of a player.	In a pathfinding problem, a state represents the <b>current location</b> of an agent in a graph.

### Q. Difference Between Uniform Cost Search (UCS) and Breadth-First Search (BFS)

Aspect	Uniform Cost Search (UCS)	Breadth-First Search (BFS)
<b>Definition</b>	A search algorithm that expands the least-cost node first.	A search algorithm that expands nodes level by level.
<b>Type of Algorithm</b>	<b>Informed Search</b> (uses path cost).	<b>Uninformed Search</b> (no cost consideration).
<b>Expansion Strategy</b>	Expands the node with the <b>lowest total path cost (g(n))</b> .	Expands all nodes at the <b>current depth</b> before moving to the next level.
<b>Uses a Cost Function?</b>	Yes, it considers the <b>cumulative cost (g(n))</b> from the start node.	No, it treats all edge costs as <b>equal</b> (assumes unit cost).
<b>Queue Type (Data Structure)</b>	<b>Priority Queue</b> (sorted by path cost).	<b>FIFO Queue</b> (First In, First Out).
<b>Optimality</b>	<b>Yes</b> , UCS finds the optimal path when costs are positive.	<b>Yes</b> , BFS finds the optimal path <b>only if all edges have the same cost</b> .
<b>Completeness</b>	<b>Yes</b> , UCS is complete if costs are non-negative.	<b>Yes</b> , BFS is complete in a finite state space.
<b>Time Complexity</b>	$O(b^{1+\text{floor}(C^*/\epsilon)})$	$O(b^d)$
<b>Space Complexity</b>	$O(b^{1+\text{floor}(C^*/\epsilon)})$	$O(b^d)$
<b>When to Use?</b>	When <b>path costs vary</b> and we need the <b>least-cost solution</b> .	When <b>all edge costs are equal</b> , and we need the <b>shortest path in terms of steps</b> .
<b>Example Use Cases</b>	Finding the <b>cheapest</b> flight between two cities, shortest path in a weighted graph.	Solving <b>mazes, shortest path problems with uniform cost</b> (e.g., unweighted graphs).

## Q. Define evaluation function or heuristic function to solve an informed search problem

### 1. Evaluation Function ( $f(n)$ ):

- In **informed search**, the **evaluation function** determines the desirability of expanding a node.
- It guides the search by assigning a numerical value to each node.
- The most common form is:  $f(n)=g(n)+h(n)$  where:
  - $g(n)$  = Cost from the start node to  $n$
  - $h(n)$  = Heuristic estimate of the cost from  $n$  to the goal.

### 2. Heuristic Function ( $h(n)$ ):

- A **heuristic function** is an approximation of the remaining cost to the goal.
- It is problem-specific and helps the algorithm prioritize nodes.
- A good heuristic function is **efficient to compute and leads the search efficiently towards the goal**.

### *Example: A Search in a Grid (Manhattan Distance Heuristic)\**

Consider a **grid-based pathfinding problem**, where you must move from **start (2,2)** to **goal (6,6)** using up, down, left, or right moves.

#### **Heuristic Calculation (Manhattan Distance)**

One common heuristic for grid-based search is the **Manhattan Distance**, given by:

$$h(n) = |x_{\text{goal}} - x_n| + |y_{\text{goal}} - y_n|$$

For a node at (2,2) with a goal at (6,6):

$$h(2,2) = |6-2| + |6-2| = 4+4=8$$

If a move costs **1 unit**, this heuristic gives a reasonable estimate of how far we are from the goal.

### **Choosing a Good Heuristic**

A heuristic should be:

1. **Admissible** → Never overestimates the actual cost.
2. **Consistent (Monotonicity Condition)** → If moving from node A to B incurs cost  $c$   
 $h(A) \leq h(B) + c$
3. **Computationally Efficient** → Should be quick to compute.

**Q. Design the heuristic functions for the 8 puzzle problem and show that the heuristic functions are admissible**

The **8-puzzle problem** consists of a **3×3 grid** with 8 numbered tiles and one empty space. The goal is to reach a specific arrangement from a given initial configuration by sliding tiles into the empty space.

**Two Common Heuristic Functions**

**1. Misplaced Tiles Heuristic ( $h_1(n)$ )**

- Counts the number of tiles **not in their goal position**.
- Example:

$$h_1(n) = \sum_{i=1}^8 \mathbb{I}(\text{tile}_i \neq \text{goal\_position}_i)$$

- $\mathbb{I}$  is an indicator function (1 if true, 0 otherwise).

**2. Manhattan Distance Heuristic ( $h_2(n)$ )**

- Computes the sum of **horizontal and vertical** moves needed for each tile to reach its goal position.
- Given by:

$$h_2(n) = \sum_{i=1}^8 (|x_i - x_{\text{goal}}| + |y_i - y_{\text{goal}}|)$$

**Proof of Admissibility**

A heuristic is **admissible** if it **never overestimates** the actual cost to the goal.

**1. Misplaced Tiles Heuristic ( $h_1(n)$ )**

- Each misplaced tile requires at least **one move** to reach the correct position.
- Since each move costs **exactly 1**, the heuristic never **overestimates** the number of moves.
- **Thus,  $h_1(n)$  is admissible.**

**2. Manhattan Distance Heuristic ( $h_2(n)$ )**

- Each tile must move at least **as many steps as Manhattan Distance suggests**.
- **No tile can reach its goal in fewer moves than its Manhattan Distance.**
- Since **tile swaps are not allowed**,  $h_2(n)$  is a **lower bound** on the actual cost.
- **Thus,  $h_2(n)$  is admissible.**

**Which Heuristic is Better?**

- **$h_1(n)$  (Misplaced Tiles) → Simpler but less accurate.**
- **$h_2(n)$  (Manhattan Distance) → More informative and generally performs better in A\* search.**

## Q. Difference Between A\* and AO\* Search Algorithm

Aspect	A* Search Algorithm	AO* Search Algorithm
<b>Type of Search</b>	Finds the <b>shortest path</b> in a state-space graph.	Finds the <b>optimal solution</b> in an AND-OR graph.
<b>Search Space</b>	Works in <b>state-space graphs or trees</b> .	Works in <b>AND-OR graphs</b> , where nodes represent <b>decisions and subproblems</b> .
<b>Graph Type</b>	Uses <b>single-path search</b> .	Uses <b>graph search with AND-OR nodes</b> (useful in problem decomposition).
<b>Node Type</b>	Each node represents a <b>state</b> .	Nodes can be <b>AND nodes (subproblems must be solved together)</b> or <b>OR nodes (one subproblem is sufficient to solve the problem)</b> .
<b>Expansion</b>	Expands the most promising node based on $f(n)=g(n)+h(n)$ .	Expands nodes <b>recursively</b> based on subproblem dependencies.
<b>Heuristic Function</b>	Uses a heuristic function $h(n)$ to estimate the cost from node $n$ to the goal.	Uses a heuristic function that <b>guides search in AND-OR graphs</b> , considering both subproblems and their dependencies.
<b>Cost Function</b>	Uses $f(n)=g(n)+h(n)$ , where: - $g(n)$ = cost from start to $n$ . - $h(n)$ = estimated cost from $n$ to goal.	Uses a cost function based on <b>aggregated cost of all subproblems</b> in AND-OR graphs.
<b>Purpose</b>	Used for <b>pathfinding and shortest path problems</b> .	Used in <b>hierarchical problem-solving and game trees</b> .
<b>Application Areas</b>	- Pathfinding (e.g., Google Maps, Robotics). - Game AI (minimax search). - Planning and scheduling.	- Expert systems. - Hierarchical problem-solving (e.g., medical diagnosis). - Decision-making in uncertain environments.
<b>Optimality</b>	<b>Guaranteed optimal solution</b> if $h(n)$ is <b>admissible and consistent</b> .	Finds an <b>optimal solution in an AND-OR graph</b> but depends on <b>how problems decompose</b> .
<b>Computational Efficiency</b>	Can be computationally expensive for large state spaces.	Efficient in <b>hierarchical problem-solving</b> but can be complex when AND-OR dependencies are large.

Q.