

## ARTIFICIAL INTELLIGENCE PyQ ANSWERS

**Explain the role of discount factor in RL, considering  $\gamma = 0, 1$  and varies b/w 0.2 to 0.8**

The **discount factor** ( $\gamma$ ) in Reinforcement Learning (RL) plays a crucial role in determining how much future rewards contribute to the agent's decision-making. It is a value between **0 and 1** that balances **immediate vs. future rewards** in the **return (cumulative reward)** calculation.

### 1. Mathematical Role of Discount Factor

The **return** at time step  $t$ :

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$$

where:

- $R_{t+1}$  is the reward at time  $t + 1$ ,
- $\gamma$  determines how much the agent values future rewards.

A **higher  $\gamma$**  makes the agent **more future-focused**, while a **lower  $\gamma$**  makes it **short-sighted**.

### 2. Effects of Different Values of $\gamma$

#### (i) When $\gamma = 0$

- The agent **only considers immediate rewards** and **completely ignores future rewards**.
- The return simplifies to  $G_t = R_{t+1}$ , meaning it behaves in a **greedy manner**, maximizing only the next reward.
- Useful in **one-step decision-making** problems where only the next action matters (e.g., reflex-based tasks).

#### (ii) When $\gamma = 1$

- The agent **considers the entire future rewards without discounting**.
- It aims for the **longest-term reward maximization**, making it **highly strategic**.
- However, in **infinite-horizon problems**, the return may **not converge**, making it computationally unstable.

#### (iii) When $\gamma$ varies (0.2 to 0.8)

- $\gamma = 0.2 \rightarrow$  The agent values **immediate rewards much more** and slightly considers future rewards.
- $\gamma = 0.5 \rightarrow$  The agent balances **short-term and long-term rewards**.
- $\gamma = 0.8 \rightarrow$  The agent **strongly considers future rewards**, optimizing for a **longer horizon** while still discounting somewhat.

### 3. Practical Considerations

- **Small  $\gamma$  (e.g., 0.2 - 0.4):** Good for **short-term tasks** like **robotic arms**, where immediate feedback is crucial.
- **Medium  $\gamma$  (e.g., 0.5 - 0.7):** Balanced strategy for **episodic tasks** like **board games**, where future outcomes matter but immediate actions are still important.

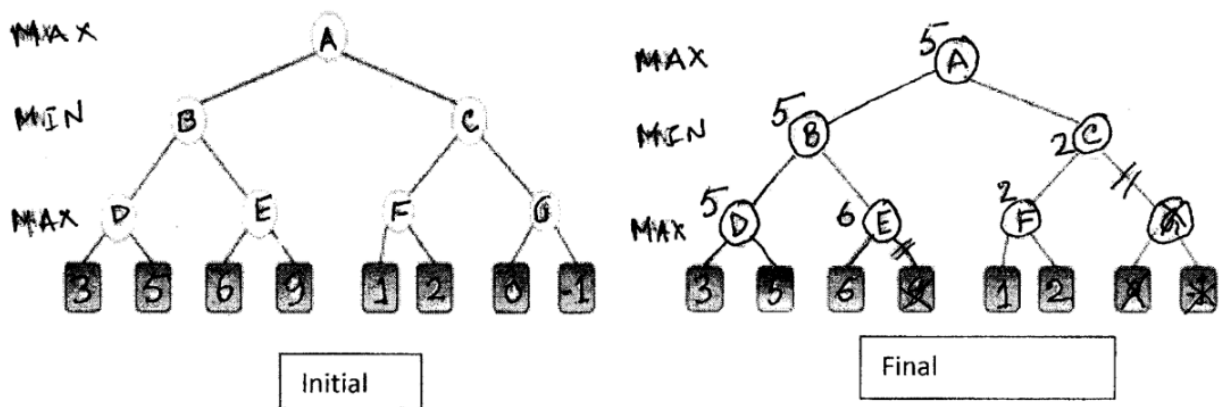
- **High  $\gamma$  (e.g., 0.8 - 0.99):** Preferred for **long-horizon tasks** like **autonomous driving** or **stock trading**, where long-term success is critical.

#### 4. Choosing $\gamma$ in Practice


- If  $\gamma$  is **too low**, the agent acts **short-sighted** and may **miss optimal strategies**.
- If  $\gamma$  is **too high**, the agent may **struggle with long-term credit assignment** and may not learn efficiently.

Q.


(i) Initial and Final game trees are given below. Explain how the final game tree is achieved using appropriate algorithm.



Ans:




(ii) At E



$$\alpha = 6; \beta = 5$$


As  $\alpha \geq \beta$ , so prune right subtree

(iii) At F



Both 1 and 2 are less than 5

(iv) At C,



As  $\alpha \geq \beta$ , Right subtree of C is pruned.

## Q. Write the properties of MINIMAX game search algorithm

The **Minimax algorithm** is used in **two-player, zero-sum games** like Chess, Tic-Tac-Toe, and Checkers. It systematically explores possible moves, assuming both players play optimally.

### 1. Completeness

**Minimax is complete** if the search tree is finite.

It guarantees finding a solution if a terminal state exists.

**Example:** In **Tic-Tac-Toe**, Minimax explores all possible moves, ensuring it finds a winning or drawing strategy.

### 2. Optimality

**Minimax is optimal** if both players play perfectly.

It assumes the opponent plays optimally and chooses the best possible move to minimize the worst-case loss.

**Example:** In **Chess**, Minimax ensures the best possible outcome based on available information.

### 3. Time Complexity

**Minimax has exponential time complexity** of  $O(b^d)$ , where:

- $b$  = branching factor (average number of moves per turn).
- $d$  = depth of the game tree.

**Example:** In Chess ( $b \approx 35$ ,  $d \approx 100$ ), Minimax becomes infeasible without optimizations like Alpha-Beta Pruning.

### 4. Space Complexity

**Depends on the implementation:**

**DFS-based Minimax**  $\rightarrow O(d)$  (depth-first, stores only one path at a time).

**BFS-based Minimax**  $\rightarrow O(b^d)$  (breadth-first, stores the entire tree).

**Example:**

- **Tic-Tac-Toe** (small tree)  $\rightarrow$  Can store the full tree.
- **Chess** (huge tree)  $\rightarrow$  Uses depth-limited search with Alpha-Beta pruning.

### 5. Deterministic & Zero-Sum

**Minimax works in deterministic, zero-sum games.**

- **Deterministic:** No randomness; every move leads to a known state.

- **Zero-Sum:** One player's gain is another's loss.

**Example:**

- **Applicable**  $\rightarrow$  Chess, Tic-Tac-Toe (Fixed moves, no randomness).
- **Not applicable**  $\rightarrow$  Poker (Random cards, bluffing).

### 6. Limited by Depth and Pruning

**Minimax is inefficient for large trees** but can be improved using:

- **Depth-limited Minimax** – Stops at a fixed depth ( $d$ ).

- **Alpha-Beta Pruning** – Reduces explored nodes, improving efficiency.

**Example:** In Chess, Alpha-Beta Pruning reduces  $b^d$  complexity to  $b^{(d/2)}$ , making deeper searches feasible.



## Q. When do you apply Alpha-Beta Pruning in the Minimax Tree?

**Alpha-Beta Pruning** is applied when we can **avoid evaluating parts of the Minimax tree** that won't affect the final decision. It helps **reduce the number of nodes explored**, making Minimax faster **without changing the result**.

### 1. Pruning Condition

**Prune a branch if we find that it cannot influence the final decision.**

- $\alpha$  (alpha) → Best value for MAX (maximizing player) so far
- $\beta$  (beta) → Best value for MIN (minimizing player) so far
- If  $\alpha \geq \beta$ , further exploration is **useless**, and we prune that branch.

### 2. When to Apply Alpha-Beta Pruning?

**Pruning occurs in two cases:**

#### 1. Beta Cutoff ( $\beta \leq \alpha$ ) in the Maximizing Level

- If a MAX node finds a move with a value  $\geq \beta$ , further children **are ignored**.

#### 2. Alpha Cutoff ( $\alpha \geq \beta$ ) in the Minimizing Level

- If a MIN node finds a move with a value  $\leq \alpha$ , further children **are ignored**.

### 3. Benefits of Alpha-Beta Pruning

- Reduces nodes explored from  $O(b^d)$  to  $O(b^{(d/2)})$  → Much faster!
- Works best when the tree is sorted (Best moves first).
- No change in the final Minimax decision.

### 4. When to Avoid Alpha-Beta Pruning?

- If the tree is unstructured/random, pruning may not help much.
- Not useful in non-deterministic games (like Poker, where chance affects outcomes).
- Sorting moves before searching increases efficiency but adds extra cost.

## Q. What is the purpose of a Belief Network ?

A **Belief Network**, also known as a **Bayesian Network (BN)**, is a **probabilistic graphical model** that represents **dependencies** among random variables using **Directed Acyclic Graphs (DAGs)**. It is used for **reasoning under uncertainty** in AI.

### Purpose of a Belief Network

#### 1. Probabilistic Reasoning

- Helps **infer hidden (unknown) variables** based on known evidence.
- Computes the **probability of events** occurring.

**Example:**

- **Medical Diagnosis:** If a patient has a fever, what is the probability they have the flu?
- **Spam Detection:** Given features like sender and keywords, what is the probability an email is spam?

#### 2. Handling Uncertainty in AI

- Real-world AI applications involve uncertainty (e.g., noisy data, incomplete info).
- Bayesian Networks model relationships probabilistically, unlike deterministic logic.

**Example:**

- A **robotic system** must decide if an object in front is a wall or a door based on noisy sensor data.

### 3. Causal Relationship Representation

- Unlike simple probability models, BNs **represent cause-and-effect relationships**.
- Helps AI **predict outcomes** when conditions change.

**Example:**

- **Traffic Prediction:** If it rains, what is the probability of a traffic jam?
  - Rain → Slippery Roads → More Accidents → Traffic Jam

#### 4. Decision Making

- Used in **decision-making systems** to evaluate different actions and their probabilities.
- Supports **decision trees, reinforcement learning, and AI agents**.

**Example:**

- **Self-driving cars:** Given current road conditions and pedestrian movement, what is the best driving action to take?

#### 5. Learning from Data

- Bayesian Networks can be **built from data** using **Bayesian inference**.
- Allows AI to **learn probabilistic dependencies** and improve over time.

**Example:**

- AI learns **which symptoms are highly correlated** with specific diseases by analyzing **medical datasets**.

### Why is Probabilistic Reasoning Needed in AI?

#### What is Probabilistic Reasoning?

Probabilistic reasoning in AI deals with **uncertainty** by assigning probabilities to different outcomes. Instead of making rigid, deterministic decisions, AI can **infer and predict outcomes based on likelihoods**. It is used in **Bayesian networks, Hidden Markov Models, Decision Trees, and Reinforcement Learning**.

### Why Do We Need Probabilistic Reasoning in AI?

#### 1. Handling Uncertainty

- **Real-world data is incomplete, noisy, or ambiguous**—probabilistic reasoning allows AI to make the **best possible decision** even when full information isn't available.
- AI must **infer missing details** instead of assuming absolute truths.

**Example:**

- A self-driving car detects a blurry object ahead. Is it **a pedestrian or just a shadow?**
- Using probabilities, the AI can determine the **most likely scenario** and react accordingly.

#### 2. Making Rational Decisions

- AI applications like **medical diagnosis, stock prediction, and robotics** require **decision-making under uncertainty**.
- Probabilistic models help AI **weigh different possibilities** and choose the **most rational action**.

**Example:**

- In **medical diagnosis**, if a patient has symptoms A, B, and C, what is the probability of **Disease X vs. Disease Y?**
- AI computes probabilities and recommends **the most likely diagnosis**.

### 3. Learning from Data

- AI can learn **patterns and trends from data** using **probability distributions**.
- Unlike rule-based systems, probabilistic models can **adapt and update** based on new information.

**Example:**

- A **spam filter** assigns a probability score based on **words, sender, and email history** to decide if a message is spam or not.
- If a user marks an email as spam, the AI **updates its probability model** to improve future predictions.

#### **4. Modeling Cause-and-Effect**

- Probabilistic reasoning allows AI to **understand causal relationships** rather than just correlations.
- Helps in **predictive modeling** where past events influence future outcomes.

**Example:**

- **Traffic Prediction System:**
  - If it **rains**, the probability of **traffic congestion** increases.
  - If it **rains and there's an accident**, congestion probability is **even higher**.

#### **5. Optimizing AI Performance**

- AI models like **Hidden Markov Models (HMMs), Bayesian Networks, and Reinforcement Learning** use probability to **balance exploration vs. exploitation**.
- This improves AI's ability to **adapt dynamically**.

**Example:**

- **Reinforcement Learning in Games**
  - AI **chooses moves based on the probability** of winning.
  - Over time, it learns which actions are more **rewarding** and adjusts its strategy.

#### **Applications of Probabilistic Reasoning in AI**

- **Robotics** – Navigate uncertain environments.
- **Natural Language Processing (NLP)** – Understand speech and text ambiguities.
- **Medical Diagnosis** – Predict diseases based on symptoms.
- **Fraud Detection** – Identify suspicious transactions using probability.
- **Self-Driving Cars** – Make safe driving decisions under uncertainty.
- **Weather Forecasting** – Predict rain, storms, or temperature changes.

## Difference Between Games and Search Problems in AI

Aspect	Games in AI	Search Problems in AI
Definition	Games involve <b>two or more agents</b> competing to achieve a goal, where each agent's actions affect the others.	Search problems involve finding a <b>sequence of actions</b> that leads to a desired goal state.
Number of Agents	<b>Multi-agent</b> environment with competing entities.	<b>Single-agent</b> environment, solving a problem independently.
Nature of Environment	<b>Adversarial</b> , as agents have <b>conflicting</b> objectives (e.g., one wins, the other loses).	<b>Non-adversarial</b> , as there is no competition, only finding an optimal solution.
Objective	To <b>maximize an agent's utility</b> while minimizing the opponent's success.	To find the <b>best or optimal path</b> from an initial state to the goal state.
Decision Process	Agents make <b>strategic decisions</b> based on the opponent's possible moves.	The search algorithm explores possible paths systematically to find a solution.
Types of Problems	Chess, Tic-Tac-Toe, Go, Poker, AlphaGo.	Route Planning, Puzzle Solving, Pathfinding, AI Planning.
Evaluation	Uses a <b>utility function</b> or evaluation function to decide the best move.	Uses <b>heuristics, cost functions, and goal tests</b> to evaluate paths.
Complexity	Often <b>more complex</b> due to the need to predict an opponent's moves (e.g., exponential growth in possibilities).	Complexity depends on the <b>state space</b> and branching factor but is usually <b>deterministic</b> .
Algorithms Used	<b>Minimax, Alpha-Beta Pruning, Monte Carlo Tree Search (MCTS).</b>	<b>A*, BFS, DFS, Dijkstra's Algorithm, Greedy Best-First Search.</b>
Example of States	In Chess, a state represents <b>board positions of all pieces</b> and the turn of a player.	In a pathfinding problem, a state represents the <b>current location</b> of an agent in a graph.

## Exploration vs. Exploitation Dilemma in AI

### 1. Definition

The **exploration vs. exploitation dilemma** is a fundamental trade-off in **decision-making AI systems**, especially in **reinforcement learning (RL)** and **multi-armed bandit (MAB)** problems.

- **Exploration:** The AI **tries new actions** to discover **better long-term rewards**.
- **Exploitation:** The AI **chooses the best-known action** to **maximize immediate reward**.

The challenge is to **balance both**—too much exploration wastes time on suboptimal actions, while too much exploitation might prevent discovering better alternatives.

### 2. Why is this Important?

- If an AI **explores too much**, it wastes time trying bad choices.
- If an AI **exploits too much**, it might get stuck in a **local optimum** and miss better options.



**Example:**

Imagine an **AI-powered investment system** that picks stocks:

- **Exploration:** Invests in new stocks to **test their profitability**.
- **Exploitation:** Continues investing in **the best-performing stocks so far**.  
Balancing both helps the AI find the **most profitable portfolio** over time.

**3. Consequences of More Exploration or More Exploitation**

Scenario	Effect
<b>Too much exploration</b>	The AI keeps trying <b>new, uncertain choices</b> , leading to <b>inconsistent rewards</b> and slow learning.
<b>Too much exploitation</b>	The AI sticks to <b>known best choices</b> , missing <b>better opportunities</b> and getting trapped in <b>suboptimal strategies</b> .

**Example:** In a **self-driving car**, choosing between two routes:

- **Exploration:** Tries different roads to find a faster route.
- **Exploitation:** Always takes the current best-known route.
- **Risk:** If it never explores, it might never discover a **better traffic-free road**.

**4. Strategies to Balance Exploration and Exploitation**

Strategy	Description
<b><math>\epsilon</math>-Greedy Algorithm</b>	AI selects a random action <b><math>\epsilon\%</math> of the time</b> (exploration) and the best-known action <b><math>(1-\epsilon)\%</math> of the time</b> (exploitation).

**Q. How many number of nodes are generated in Depth Limited Search and Iterative Depth Search Algorithms considering depth  $d=4$  and branching factor  $b=12$**

$$\begin{aligned}
 b &= 12, d = 4 \\
 N_{DLS} &= b^0 + b^1 + \dots + b^d \\
 &= 12^0 + 12^1 + \dots + 12^4 \\
 &= 1 + 12 + 144 + 1728 + 20736 \\
 &= 22621 \\
 N_{IDS} &= (d+1)b^0 + db^1 + (d-1)b^2 + \dots + 2b^{d-1} + b^d \\
 &= 5 \times 1 + 4 \times 12^1 + 3 \times 12^2 + 2 \times 12^3 + 1 \times 12^4 \\
 &= 5 + 48 + 432 + 3456 + 20736 \\
 &= 24677
 \end{aligned}$$

### Q. Difference Between Uniform Cost Search (UCS) and Breadth-First Search (BFS)

Aspect	Uniform Cost Search (UCS)	Breadth-First Search (BFS)
<b>Definition</b>	A search algorithm that expands the least-cost node first.	A search algorithm that expands nodes level by level.
<b>Type of Algorithm</b>	<b>Informed Search</b> (uses path cost).	<b>Uninformed Search</b> (no cost consideration).
<b>Expansion Strategy</b>	Expands the node with the <b>lowest total path cost (g(n))</b> .	Expands all nodes at the <b>current depth</b> before moving to the next level.
<b>Uses a Cost Function?</b>	Yes, it considers the <b>cumulative cost (g(n))</b> from the start node.	No, it treats all edge costs as <b>equal</b> (assumes unit cost).
<b>Queue Type (Data Structure)</b>	<b>Priority Queue</b> (sorted by path cost).	<b>FIFO Queue</b> (First In, First Out).
<b>Optimality</b>	<b>Yes</b> , UCS finds the optimal path when costs are positive.	<b>Yes</b> , BFS finds the optimal path <b>only if all edges have the same cost</b> .
<b>Completeness</b>	<b>Yes</b> , UCS is complete if costs are non-negative.	<b>Yes</b> , BFS is complete in a finite state space.
<b>Time Complexity</b>	$O(b^{1+\text{floor}(C^*/\epsilon)})$	$O(b^d)$
<b>Space Complexity</b>	$O(b^{1+\text{floor}(C^*/\epsilon)})$	$O(b^d)$
<b>When to Use?</b>	When <b>path costs vary</b> and we need the <b>least-cost solution</b> .	When <b>all edge costs are equal</b> , and we need the <b>shortest path in terms of steps</b> .
<b>Example Use Cases</b>	Finding the <b>cheapest</b> flight between two cities, shortest path in a weighted graph.	Solving <b>mazes, shortest path problems with uniform cost</b> (e.g., unweighted graphs).

## Q. Define evaluation function or heuristic function to solve an informed search problem

### 1. Evaluation Function ( $f(n)$ ):

- In **informed search**, the **evaluation function** determines the desirability of expanding a node.
- It guides the search by assigning a numerical value to each node.
- The most common form is:  $f(n)=g(n)+h(n)$  where:
  - $g(n)$  = Cost from the start node to  $n$
  - $h(n)$  = Heuristic estimate of the cost from  $n$  to the goal.

### 2. Heuristic Function ( $h(n)$ ):

- A **heuristic function** is an approximation of the remaining cost to the goal.
- It is problem-specific and helps the algorithm prioritize nodes.
- A good heuristic function is **efficient to compute and leads the search efficiently towards the goal**.

### *Example: A Search in a Grid (Manhattan Distance Heuristic)\**

Consider a **grid-based pathfinding problem**, where you must move from **start (2,2)** to **goal (6,6)** using up, down, left, or right moves.

#### **Heuristic Calculation (Manhattan Distance)**

One common heuristic for grid-based search is the **Manhattan Distance**, given by:

$$h(n) = |x_{\text{goal}} - x_n| + |y_{\text{goal}} - y_n|$$

For a node at (2,2) with a goal at (6,6):

$$h(2,2) = |6-2| + |6-2| = 4+4=8$$

If a move costs **1 unit**, this heuristic gives a reasonable estimate of how far we are from the goal.

### **Choosing a Good Heuristic**

A heuristic should be:

1. **Admissible** → Never overestimates the actual cost.
2. **Consistent (Monotonicity Condition)** → If moving from node A to B incurs cost  $c$   
 $h(A) \leq h(B) + c$
3. **Computationally Efficient** → Should be quick to compute.

**Q. Design the heuristic functions for the 8 puzzle problem and show that the heuristic functions are admissible**

The **8-puzzle problem** consists of a **3×3 grid** with 8 numbered tiles and one empty space. The goal is to reach a specific arrangement from a given initial configuration by sliding tiles into the empty space.

**Two Common Heuristic Functions**

**1. Misplaced Tiles Heuristic ( $h_1(n)$ )**

- Counts the number of tiles **not in their goal position**.
- Example:

$$h_1(n) = \sum_{i=1}^8 \mathbb{I}(\text{tile}_i \neq \text{goal\_position}_i)$$

- $\mathbb{I}$  is an indicator function (1 if true, 0 otherwise).

**2. Manhattan Distance Heuristic ( $h_2(n)$ )**

- Computes the sum of **horizontal and vertical** moves needed for each tile to reach its goal position.
- Given by:

$$h_2(n) = \sum_{i=1}^8 (|x_i - x_{\text{goal}}| + |y_i - y_{\text{goal}}|)$$

**Proof of Admissibility**

A heuristic is **admissible** if it **never overestimates** the actual cost to the goal.

**1. Misplaced Tiles Heuristic ( $h_1(n)$ )**

- Each misplaced tile requires at least **one move** to reach the correct position.
- Since each move costs **exactly 1**, the heuristic never **overestimates** the number of moves.
- **Thus,  $h_1(n)$  is admissible.**

**2. Manhattan Distance Heuristic ( $h_2(n)$ )**

- Each tile must move at least **as many steps as Manhattan Distance suggests**.
- **No tile can reach its goal in fewer moves than its Manhattan Distance.**
- Since **tile swaps are not allowed**,  $h_2(n)$  is a **lower bound** on the actual cost.
- **Thus,  $h_2(n)$  is admissible.**

**Which Heuristic is Better?**

- **$h_1(n)$  (Misplaced Tiles) → Simpler but less accurate.**
- **$h_2(n)$  (Manhattan Distance) → More informative and generally performs better in A\* search.**

## Q. Difference Between A\* and AO\* Search Algorithm

Aspect	A* Search Algorithm	AO* Search Algorithm
<b>Type of Search</b>	Finds the <b>shortest path</b> in a state-space graph.	Finds the <b>optimal solution</b> in an AND-OR graph.
<b>Search Space</b>	Works in <b>state-space graphs or trees</b> .	Works in <b>AND-OR graphs</b> , where nodes represent <b>decisions and subproblems</b> .
<b>Graph Type</b>	Uses <b>single-path search</b> .	Uses <b>graph search with AND-OR nodes</b> (useful in problem decomposition).
<b>Node Type</b>	Each node represents a <b>state</b> .	Nodes can be <b>AND nodes (subproblems must be solved together)</b> or <b>OR nodes (one subproblem is sufficient to solve the problem)</b> .
<b>Expansion</b>	Expands the most promising node based on $f(n)=g(n)+h(n)$ .	Expands nodes <b>recursively</b> based on subproblem dependencies.
<b>Heuristic Function</b>	Uses a heuristic function $h(n)$ to estimate the cost from node $n$ to the goal.	Uses a heuristic function that <b>guides search in AND-OR graphs</b> , considering both subproblems and their dependencies.
<b>Cost Function</b>	Uses $f(n)=g(n)+h(n)$ , where: - $g(n)$ = cost from start to $n$ . - $h(n)$ = estimated cost from $n$ to goal.	Uses a cost function based on <b>aggregated cost of all subproblems</b> in AND-OR graphs.
<b>Purpose</b>	Used for <b>pathfinding and shortest path problems</b> .	Used in <b>hierarchical problem-solving and game trees</b> .
<b>Application Areas</b>	- Pathfinding (e.g., Google Maps, Robotics). - Game AI (minimax search). - Planning and scheduling.	- Expert systems. - Hierarchical problem-solving (e.g., medical diagnosis). - Decision-making in uncertain environments.
<b>Optimality</b>	<b>Guaranteed optimal solution</b> if $h(n)$ is <b>admissible and consistent</b> .	Finds an <b>optimal solution in an AND-OR graph</b> but depends on <b>how problems decompose</b> .
<b>Computational Efficiency</b>	Can be computationally expensive for large state spaces.	Efficient in <b>hierarchical problem-solving</b> but can be complex when AND-OR dependencies are large.